

Introduction to MongoDB

CS185C: Introduction to NoSQL Databases

Suneuy Kim

Reference

- [1] <https://docs.mongodb.com/manual/>
- [2] The Definitive Guide to MongoDB: A Complete Guide to Dealing with Big Data using MongoDB, Third Edition by David Howes, Peter Membrey, Eelco Plugge and Tim Hawkins, December 16, 2015
- [3] MongoDB: The Definitive Guide, 2nd Edition, Powerful and Scalable Data Storage by Kristina Chodorow, May 2013

Using the Right Tool for the Right Job

- One size (relational databases) does not fit all (different types of data)
- MongoDB provides rich document-oriented database that's optimized for speed and scalability
- Polygot Persistency (e.g. RDBMS for the accounting components and MongoDB for the document storage)

MongoDB Features

- WiredTiger
 - The default storage engine as of MongoDB3.2. (MMAP used to be the default.)
 - Document level locking as compared to collection level locking in MMAP
 - Compression
- Storage engine
 - A software module that a database management system uses to create, read, update data from a database.
 - Storage engine is where processes such as locking, index maintenance, and transactions occur.

JSON and MongoDB

- CSV is to store flat data, not for aggregate

Lastname, Firstname, Phone Number
Membrey, Peter, +852 1234 5678
Thielen, Wouter, +81 1234 5678

- [Q] What if someone has more than two phone numbers?

- JSON

```
{  "firstname": "Peter",
  "lastname": "Membrey",
  "numbers": [
    {
      "phone": "+852 1234 5678"
    },
    {
      "fax": "+44 1234 565 555"
    }
  ]
}
```

- JSON allows complex data structures to be represented in a simple, human-readable text format.
- MongoDB stores data in **BSON** (Binary JSON), not in JSON

Using Document-Oriented Storage - BSON

- Binary JSON developed by MongoDB
- MongoDB stores data in BSON.
- Easier to traverse and index very quickly at the cost of slightly more space than JSON (MongoDB is meant to be fast, rather than space-efficient.)
- It is easy and quick to convert BSON to a native data structure for each high level language (Python, Ruby, etc.)
- Extensions to JSON
 - Extended types for numeric data (int32 and int64) and support for binary data

Supporting Dynamic Queries

- SQL: static data and dynamic query
- CouchDB: dynamic data and static query or query in map-reduce functions
- MongoDB: dynamic data and dynamic query

Note:

- Dynamic data means schema less
- Dynamic query means you can run a query without planning for it in advance.

Index

- All documents are automatically indexed on the `_id` key – unique index
- A user defined index allows duplicates (e.g. index on last name key)
- By default, an error occurs if you try to create a unique index on a key that has duplicate values
- Indexes on embedded documents (e.g. an index on the ZIP or postal code)
- Extensive support for indexing your documents
 - Composite Indexes (e.g. an index that combines both the lastname and firstname)
 - Geospatial Indexes
 - Many more ...

Profiling Queries

- MongoDB's query planner `explain()`

```
db.media.find().explain();
```

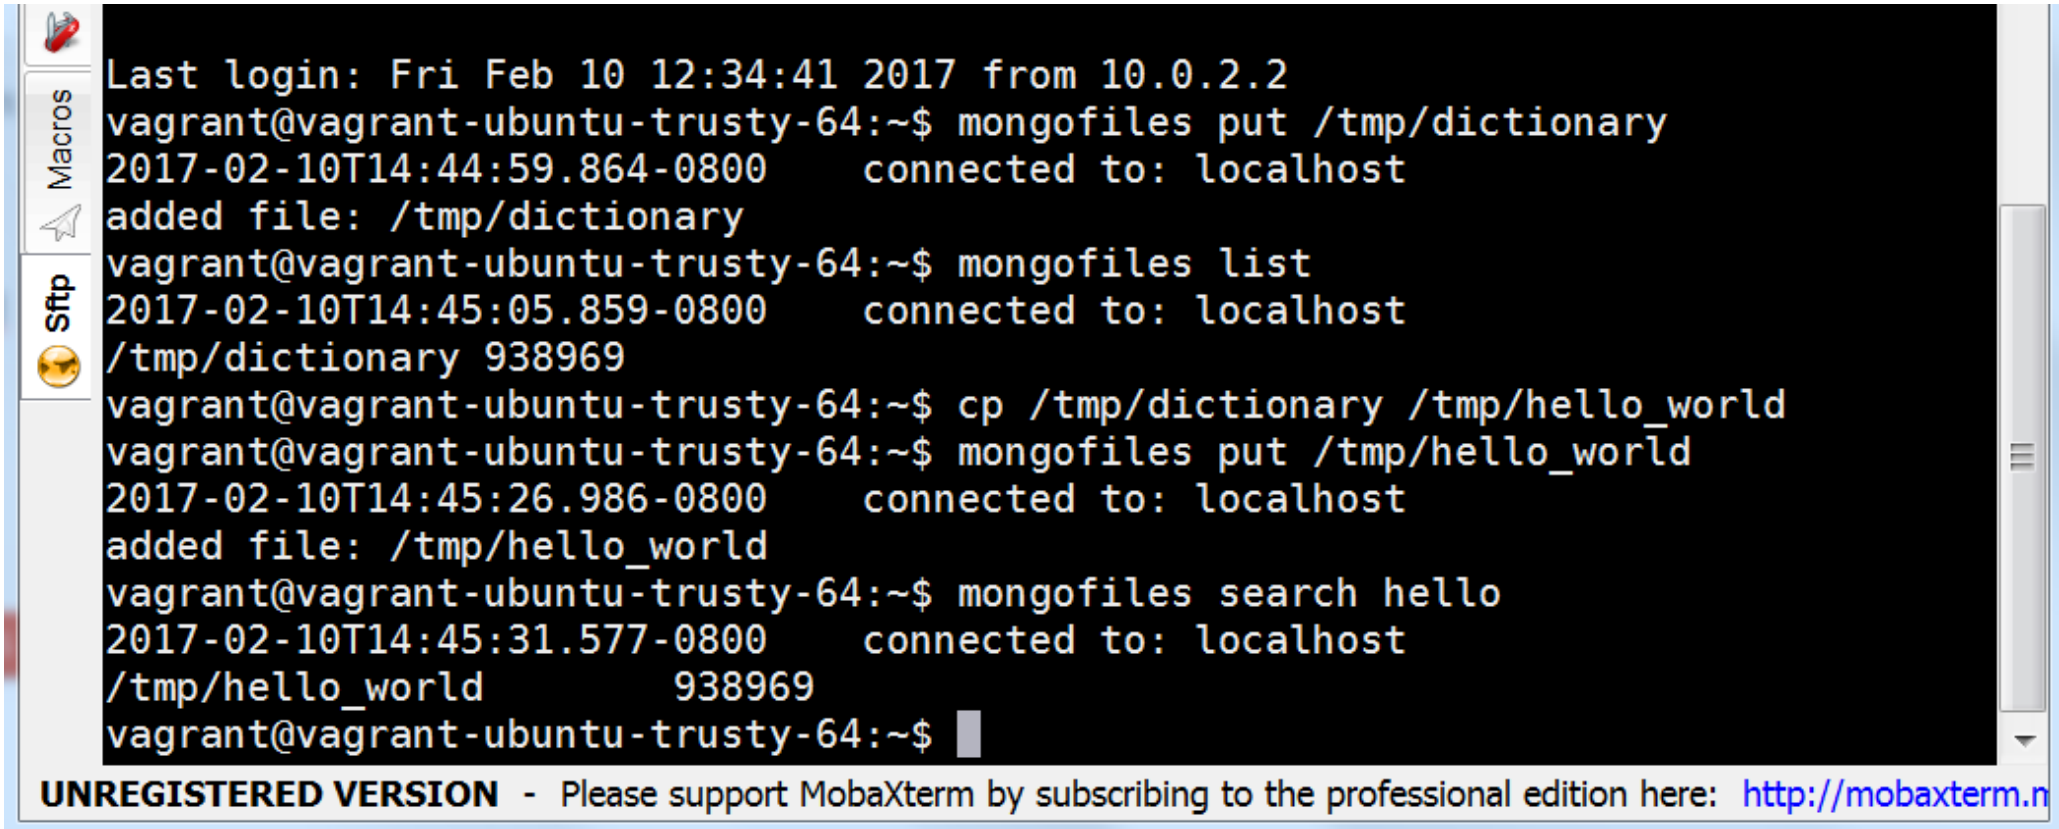
V.S.

```
db.media.find({ "_id" :  
ObjectId("58013ad941a51bd2d7599db9") }) .explain();
```

GridFS

- The maximum size of a MongoDB document in BSON: 16 MB: not enough for movie clips, high-quality audio clips, etc.
- GridFS is to store large files and yet to access parts of the file without retrieving the entire thing.
- To MongoDB, files in GridFS are just normal collections containing documents.
- GridFS consists of two collections
 - Metadata are in the `files` collection
 - Data are broken down into chunks that are stored in the `chunks` collection – easy and scalable

GridFS command line tool: mongofiles



```
Last login: Fri Feb 10 12:34:41 2017 from 10.0.2.2
vagrant@vagrant-ubuntu-trusty-64:~$ mongofiles put /tmp/dictionary
2017-02-10T14:44:59.864-0800    connected to: localhost
added file: /tmp/dictionary
vagrant@vagrant-ubuntu-trusty-64:~$ mongofiles list
2017-02-10T14:45:05.859-0800    connected to: localhost
/tmp/dictionary 938969
vagrant@vagrant-ubuntu-trusty-64:~$ cp /tmp/dictionary /tmp/hello_world
vagrant@vagrant-ubuntu-trusty-64:~$ mongofiles put /tmp/hello_world
2017-02-10T14:45:26.986-0800    connected to: localhost
added file: /tmp/hello_world
vagrant@vagrant-ubuntu-trusty-64:~$ mongofiles search hello
2017-02-10T14:45:31.577-0800    connected to: localhost
/tmp/hello_world      938969
vagrant@vagrant-ubuntu-trusty-64:~$
```

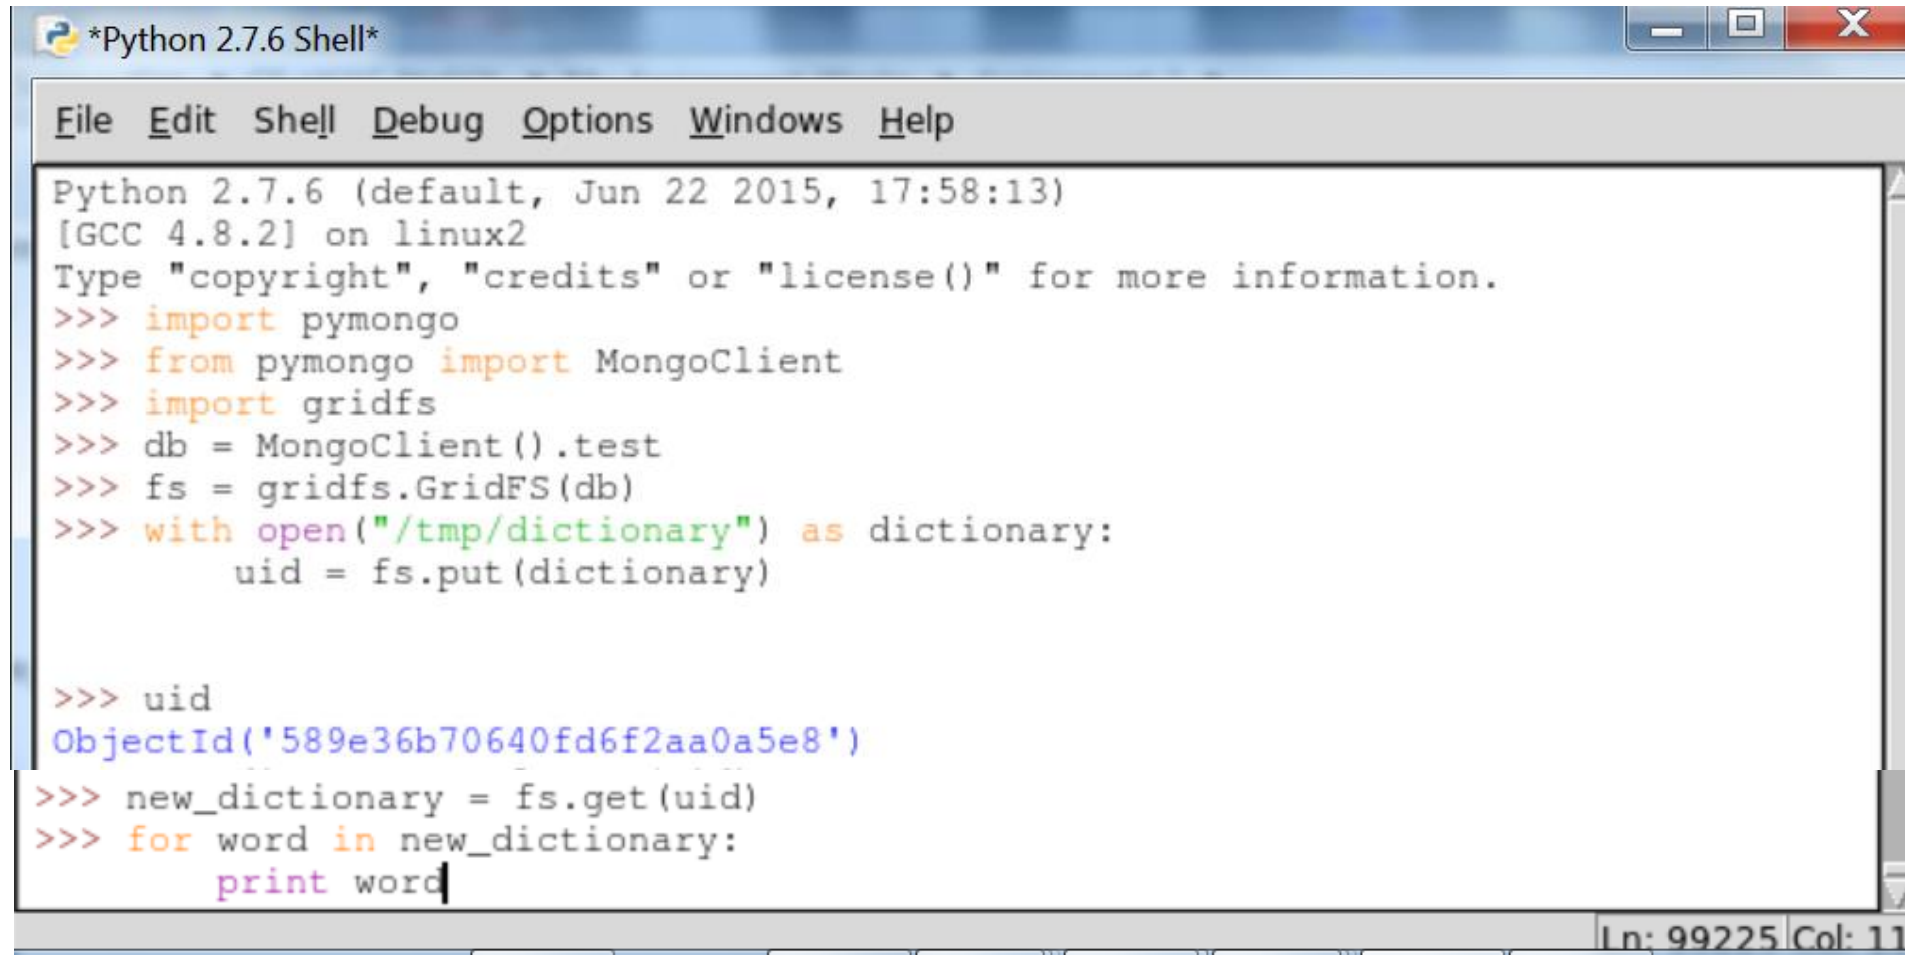
UNREGISTERED VERSION - Please support MobaXterm by subscribing to the professional edition here: <http://mobaxterm.n>

Example from [2] Chapter 5

To MongoDB, files in GridFS are just normal collections containing documents.

```
audit100
blog
comments
foo
fs.chunks
fs.files
inventory
lists
media
mediadb.media
multi
products
publisherscollection
system.indexes
texttest
texttest1
users
> db.fs.files.find()
{ "_id" : ObjectId("589e426b0640fd71967585c1"), "chunkSize" : 261120, "uploadDate" : ISODate("2017-02-10T22:44:59.911Z"), "length" : 938969, "md5" : "7e2877e5dad6e8e97b0fa43d28f2feca", "filename" : "/tmp/dictionary" }
{ "_id" : ObjectId("589e42860640fd71a7904026"), "chunkSize" : 261120, "uploadDate" : ISODate("2017-02-10T22:45:27.167Z"), "length" : 938969, "md5" : "7e2877e5dad6e8e97b0fa43d28f2feca", "filename" : "/tmp/hello_world" }
>
```

GridFS: from pymongo (python driver)

A screenshot of a Python 2.7.6 Shell window. The window has a title bar with the text '*Python 2.7.6 Shell*' and standard window controls (minimize, maximize, close). Below the title bar is a menu bar with options: File, Edit, Shell, Debug, Options, Windows, and Help. The main area of the window contains a Python REPL session. The session starts with version and environment information: 'Python 2.7.6 (default, Jun 22 2015, 17:58:13)' and '[GCC 4.8.2] on linux2'. It then prompts the user to type 'copyright', 'credits', or 'license()' for more information. The user enters several commands: 'import pymongo', 'from pymongo import MongoClient', 'import gridfs', 'db = MongoClient().test', 'fs = gridfs.GridFS(db)', and a 'with open' block to write a dictionary to GridFS. The prompt '>>>' is shown for each command. The output shows the 'uid' as an ObjectId and the retrieval of the dictionary, with words being printed. The status bar at the bottom right shows 'Ln: 99225 Col: 11'.

Example from [2] Chapter 5

Replica Sets

- A replica set has one primary server.
- The primary server handles all the write requests from the clients.
- When a write occurs, it is logged in the primary's oplog.
- The `oplog` is replicated by the secondary servers in the same replica set.
- When the primary fails, a new primary will be elected among surviving members of the replica.

Sharding

- Individual documents are self-contained - BSON
- Sharding provides horizontal scalability - additional shards can be added to increase resource capacity without any changes to your application code.
- Auto-sharding: MongoDB takes care of all the data splitting and recombination for you.

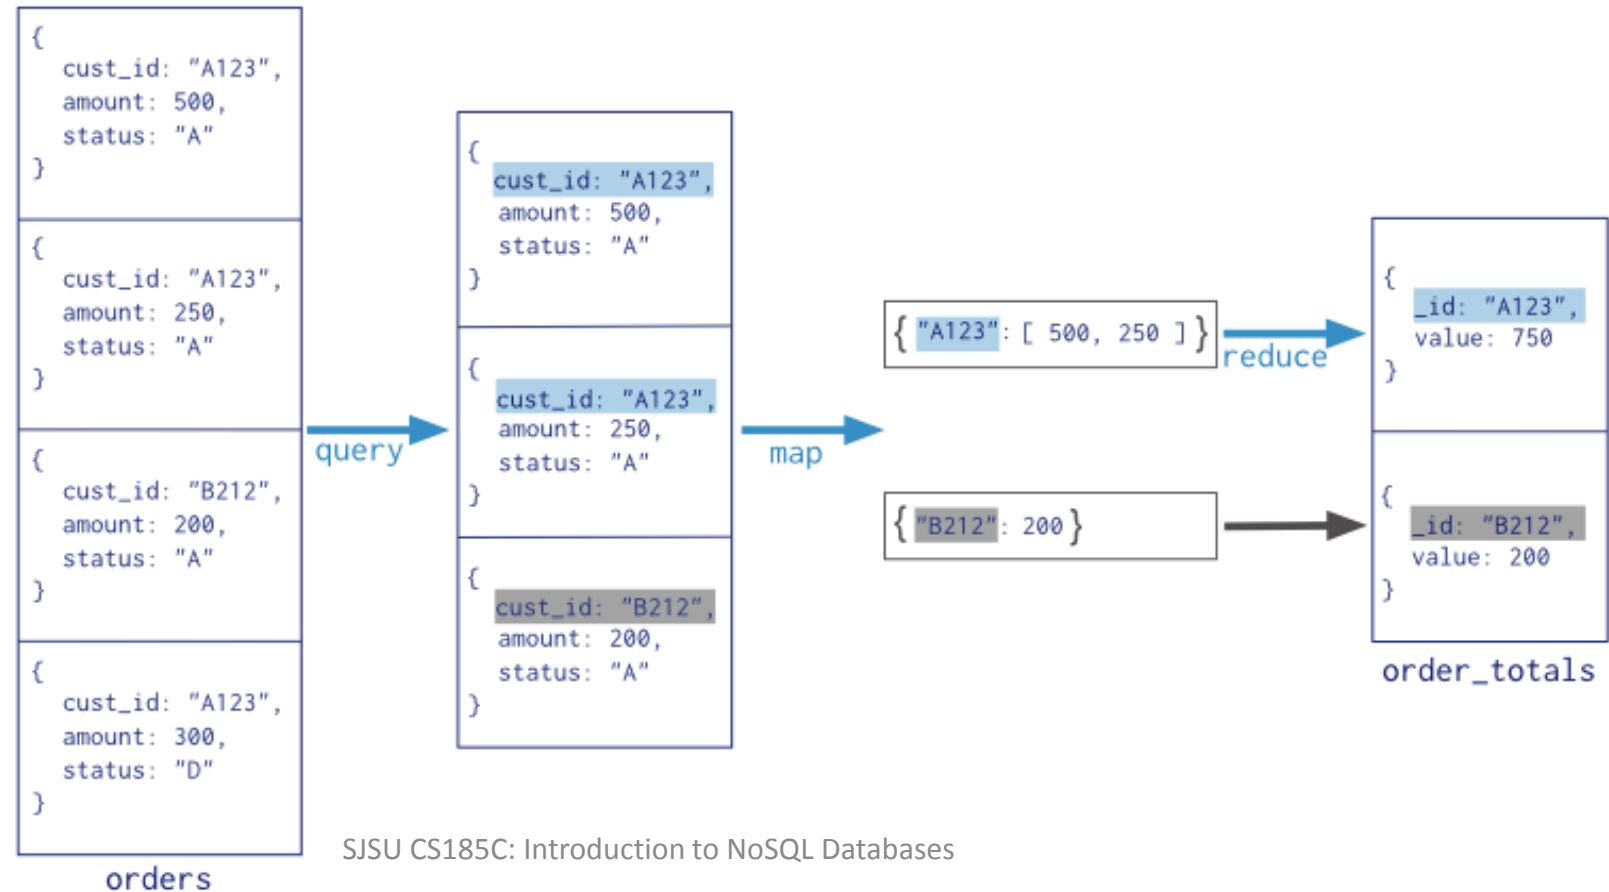
Advanced Queries

- Using Map and Reduce functions
 - Not required but provided.
 - If you would normally use GROUP BY in SQL, then the map and reduce functions may be the right tools for the job in MongoDB.
- The Aggregation Framework
 - Map-reduce can be slow.
 - Pipe line based aggregate operators implemented in C++: highly performant

MongoDB MapReduce function

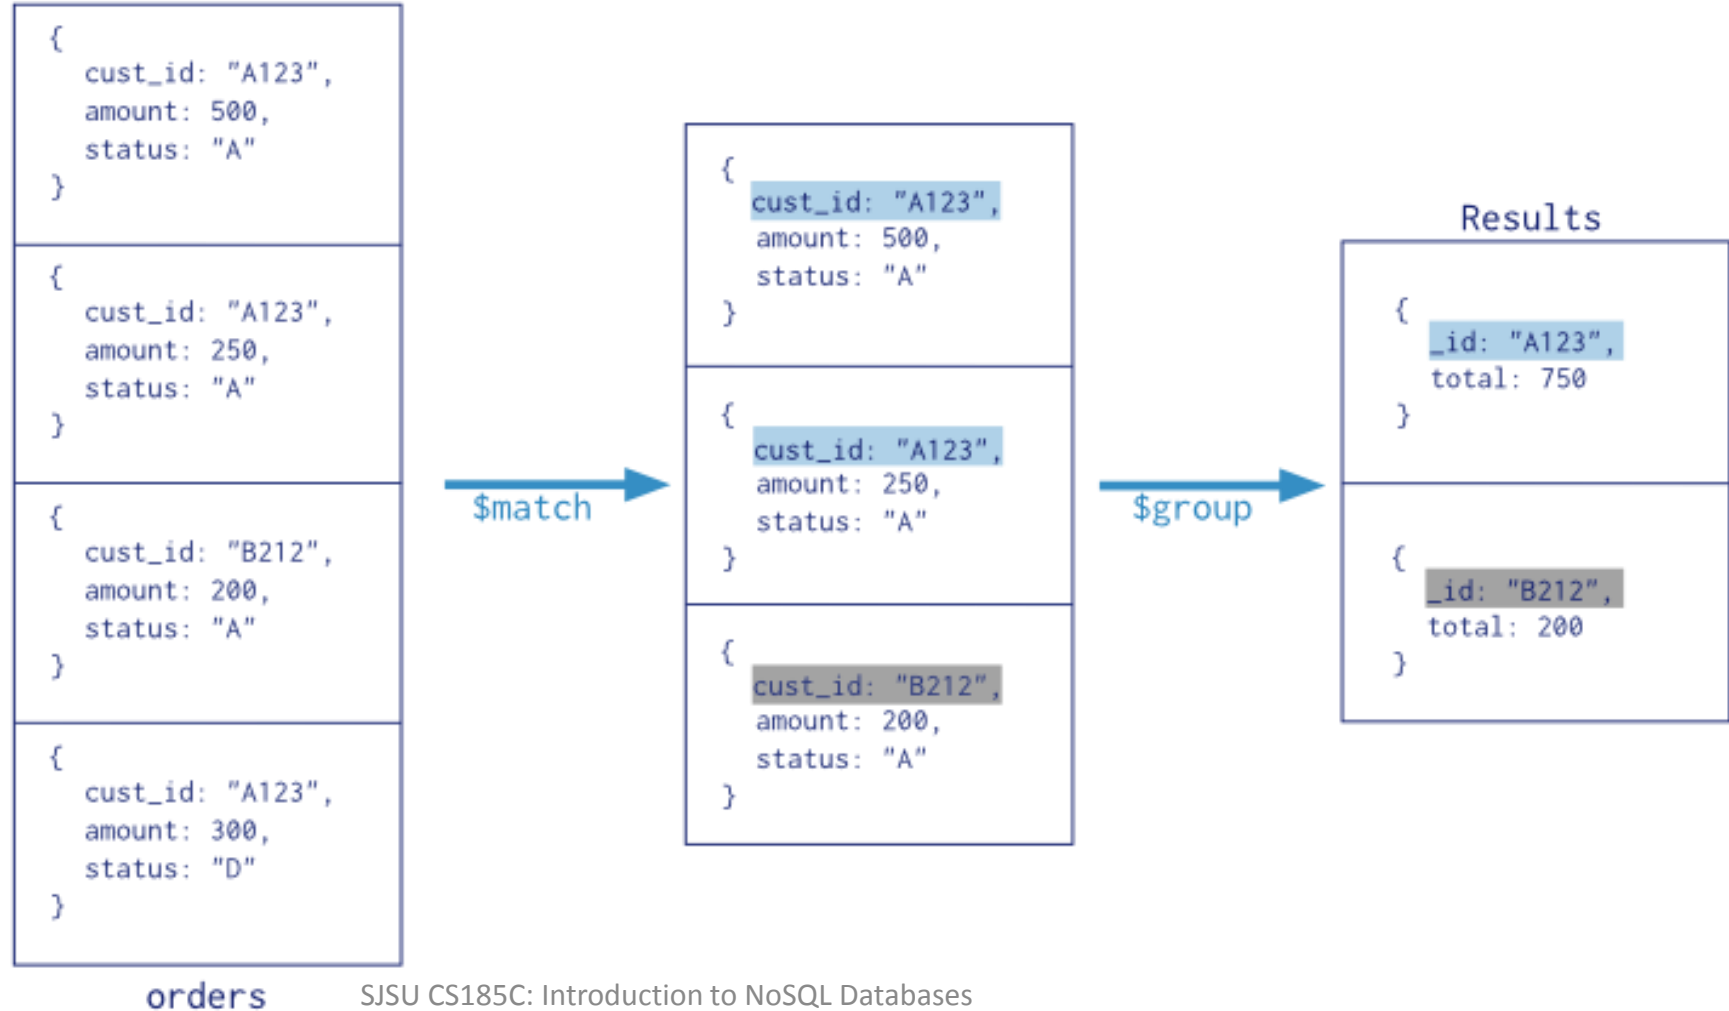
Collection
↓
db.orders.mapReduce(
 map → function() { emit(this.cust_id, this.amount); },
 reduce → function(key, values) { return Array.sum(values) },

 query → {
 output → query: { status: "A" },
 out: "order_totals"
 }
)



Pipeline based aggregation

Collection
↓
`db.orders.aggregate([`
 \$match stage → `{ $match: { status: "A" } },`
 \$group stage → `{ $group: { _id: "$cust_id", total: { $sum: "$amount" } } }`
 `]`)



Getting Started

Document

- The unit of storage in MongoDB (v.s. a row in RDBMS)
- A document an ordered set of key-value pairs.

e.g. This document `{"type" : "Book"}` contains one key-value pair consisting of a key named " type", and its value, "Book".

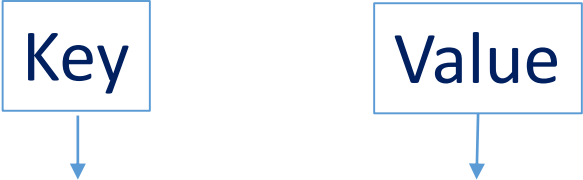
- Keys are strings
- Values can be any of the following types

String, Integer, Boolean, Double, Min/Max keys, Arrays, Timestamp, Object, Null, Symbol, Date, ObjectID, Binary data, Regular expression, Java Script code.

Document

- MongoDB is type-sensitive and case-sensitive
`{"foo", 3}` is distinct from `{"foo": "3"}` and `{"Foo": 3}`
- A MongoDB document cannot have duplicate keys
- Key/value pairs in documents are ordered.
`{"x": 1, "y": 2}` is not the same as `{"y": 2, "x": 1}`
- Not requires every document to have the same field, or that every field with the same name has the same type of value.


Document Example



```
{  
  "firstname": "Peter",  
  "lastname": "Membrey",  
  "phone_numbers": [  
    "+852 1234 5678",  
    "+44 1234 565 555"  
  ]  
}
```

`_id` field of document

- Each document has a unique identifier `_id` of which value is auto generated by default.
- `_id` type is `ObjectID` : a 12 byte unique id that can be generated independently in a **distributed sharded** environment.
- Automatically added to a new document
- The default value is a `ObjectId` BSON data type consisting of a 12-byte binary value
- Time and inc (=counter) fields are stored in Big Endian format.



Collections

- A group of similar documents.
- Dynamic Schemas: Documents within a single collection can have any number of different shapes (different keys and different types of values)
 - You do not need to predefine a structure for any of the document
 - Supports programming in a dynamic typed language such as Python or PHP
 - Still need to define collections and indexes

Collections

- It is common practice to group related types of documents together
 - Application code doesn't have to weed out irrelevant documents
 - Faster to get a list of collections than to extract a list of the types in a collection
 - Data locality
 - Efficient indexing

Collections

- Expandable collections (default) vs. capped collections
- Every collection should have a unique name
 - Should begin with a letter or an underscore (_)
 - \$ is reserved by MongoDB
 - An empty string is not allowed.
 - The null character cannot be used.
 - Cannot start with the `system.` string.
- Each collection accounts for at least two name spaces
 - One for the collection itself
 - One for the first index created in the collection.

Collections

- Sub-collection

By convention, collections are organized using name spaced sub-collection separated by .


- Example: `blog.authors`, `blog.posts`

Databases


- A database is a group of collections.
- Rule of thumbs: to store all data for a single application in the same database
- Database names will actually end up as files on your file system. (Naming restrictions exist due to this fact.)
- Reserved database names: admin, local, config
- Namespace (fully qualified collection name) = database name + collection name
e.g. cms.blog.posts

Data Model: MongoDB vs RDBMS

MongoDB Database Model



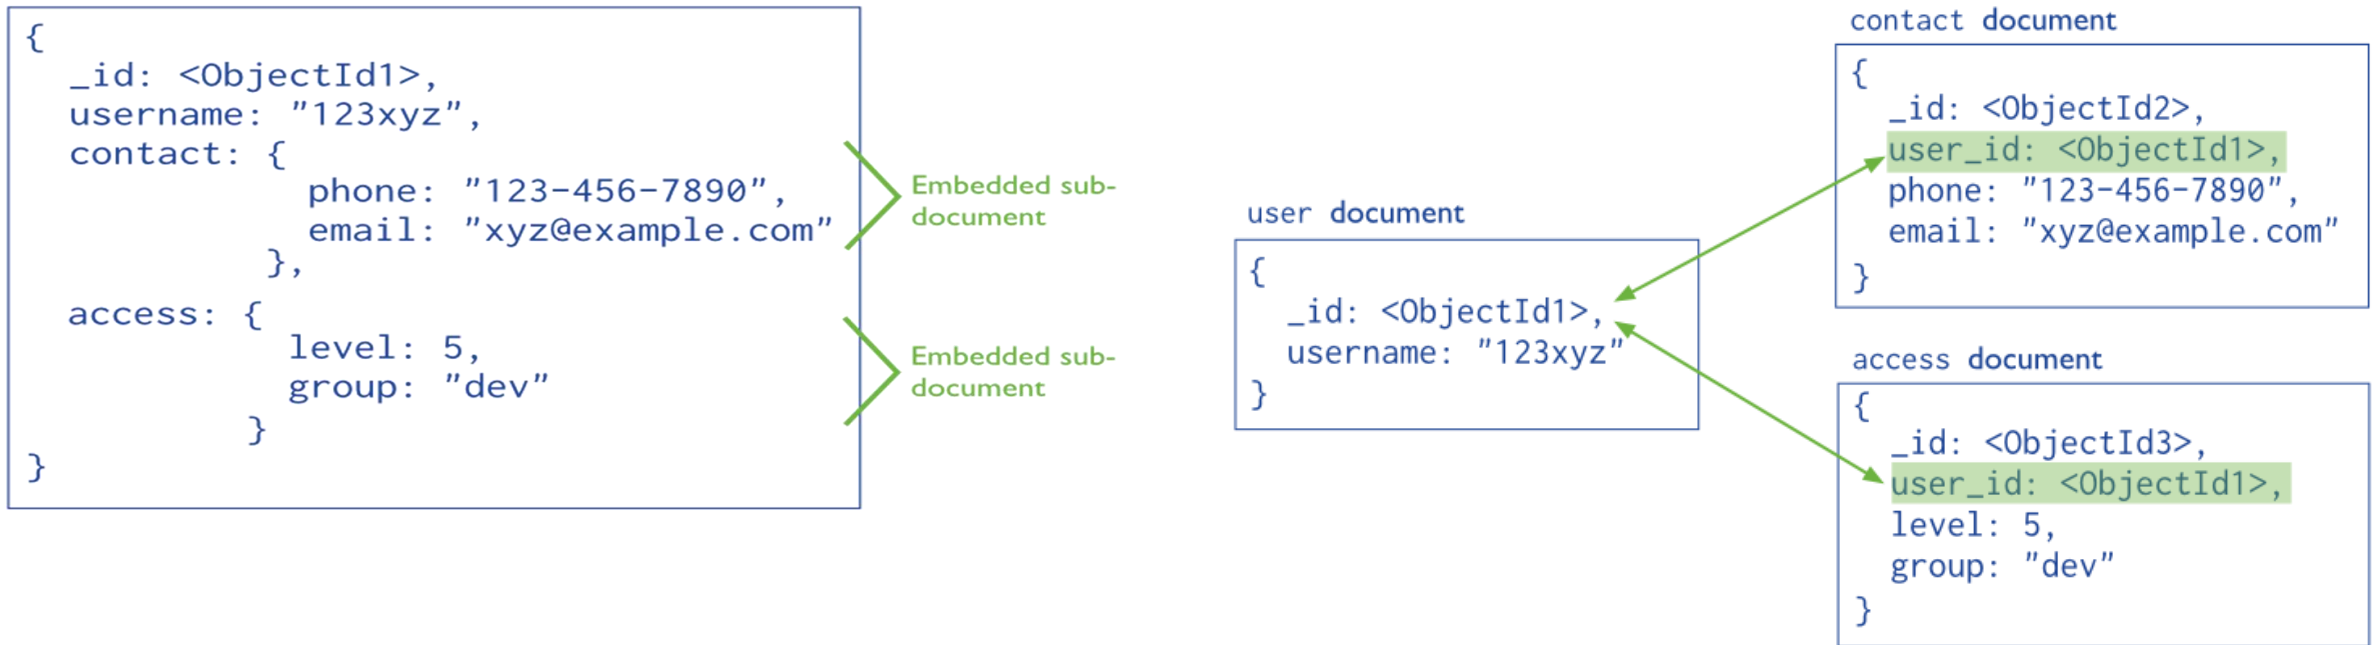
RDBMS Data Model



Embedding vs. Referencing Information in Documents

- Embedding information: you place a certain type of data (for example, an array containing more data) into the document itself.
- Referencing information: you create a reference to another document that contains that specific data (e.g. RDBMS foreign keys. Joins are needed to put information together.)
- With MongoDB, embed data whenever you can
 - To ensure that all related information is kept in one single document
 - Works much faster because the data are co-located in the disk

Embedding vs. Referencing



Every reference needs another query in the database.

Basic Data Types

- JSON's expressive capabilities are limited because the only types are null, boolean, numeric, string, array, and object.
- MongoDB adds support for a number of additional data types while keeping JSON's essential key/value pair nature.

Basic Data Types

Type	Example
null	<code>{"x" : null}</code>
boolean	<code>{"x" : true}</code>
number	<code>{"x" : 3.14} {"x" : NumberInt("3")}</code>
string	<code>{"x" : "foobar"}</code>
date	<code>{"x" : new Date()}</code>
regular expression	<code>{"x" : /foobar/i}</code>
array	<code>{"x" : ["a", "b", "c"]}</code>
embedded document	<code>{"x" : {"foo" : "bar"}}</code>
code	<code>{"x" : function() { /* ... */ }}</code>

Arrays

- Arrays can contain different data types as values
e.g.) {"things" : ["pie", 3.14]}
- MongoDB knows how to reach inside of arrays to perform operations on their contents.
e.g.) Find all documents where 3.14 is an element of the "things" array.

```
> db.foo.insert({"things":["pie",3.14]})
WriteResult({ "nInserted" : 1 })
> db.foo.find({"things":3.14})
{ "_id" : ObjectId("589f5d5ac6bf1740f89d19a1"), "things" : [ "pie", 3.14 ] }
```

Embedded Documents

- Documents can be used as the value for the key.
- Embedded "address" document

```
{  
  "name" : "John Doe",  
  "address" : {  
    "street" : "123 Park Street",  
    "city" : "Anytown",  
    "state" : "NY"  
  }  
}
```

- MongoDB is able to reach inside embedded documents to build indexes, perform queries, or make updates.

```
> db.foo.insert(customer)  
WriteResult({ "nInserted" : 1 })  
> db.foo.find({"address.state":"NY"})  
{ "_id" : ObjectId("589f671bc6bf1740f89d19a4"), "name" : "John Doe", "address" : { "street" : "123 Park Street", "city" : "Anytown", "state" : "NY" } }  
> █
```

To Setup and Deploy MongoDB

- Refer to a separate document from the course web site.

To start/stop the server

- `$sudo service mongod status`
- `$sudo service mongod start [restart|stop]`
- **The default data directory mongod uses:** `/data/db`
- **Make sure to create `/data/db` and the user has a write permission to this directory.**
- **The default port:** `27017`

mongo – Mongo Shell

- A JavaScript shell
- JavaScript equivalents to shell helpers

Helper	Equivalent
use foo	db.getSisterDB("foo")
show dbs	db.getMongo().getDBs()
show collections	db.getCollectionNames()