

CS157A: Introduction to Database Management Systems

Chapter 8: Views and Indexes

Suneuy Kim

Views

- Virtual view
 - A relations that is the result of a query over other relations.
 - Virtual views are not stored in the databases
- Materialized views are periodically constructed from the database and stored there.
 - Enable efficient access to the database
 - Are most useful in data warehousing scenarios, where frequent queries of the actual base tables can be extremely expensive.

Declaring Views

```
DROP VIEW IF EXISTS SeniorUsers;  
CREATE VIEW SeniorUsers AS  
  SELECT uID, uName, age  
  FROM User  
  WHERE age >=60;
```

Querying Views

- A view is queried as if it was a stored table.
- The query processor replaces the view by its definition to process the query.

```
SELECT *  
FROM SeniorUsers, Loan  
WHERE SeniorUsers.uID = Loan.uID and overdue = 1;
```

```
SELECT *  
FROM (SELECT uID, uName, age FROM User  
      WHERE age >=60) SU, Loan  
WHERE SU.uID = Loan.uID and overdue = 1;
```

View Removal

- `DROP VIEW SeniorUsers;`
 - Will not affect the base table `User`.
 - Can't do any query involving `SeniorUsers`.
- `DROP TABLE User;`
 - Will remove the table `User` from the database and also make all views referring the `User` table unusable.

Modifying Views

- We can't modify (insert, delete, update) a view like a table since a view is not stored.
- However, for the users who access the database only through views, there should be a way to modify views – modification to a view should be reflected on the base tables.
- Rewrite a query that modifies a view in a way that it modifies the base tables.

Example

```
DROP VIEW IF EXISTS YoungUsers;  
CREATE VIEW YoungUsers AS  
  SELECT uID, age  
  FROM User  
  WHERE age < 19;
```

Insert Into YoungUsers
Values (77, 17);


How can we translate this query to modify the
database ?

Insert into User Values (77, ?, 17); // so many
different ways !

YoungUsers

VIEW

uID	age
12	17
24	10



Modifying Views

- Correctness of translation can be achieved but ambiguity issue should be resolved and it is not straightforward for some cases.

Two main approaches of translating modification to a view

1. INSTEAD-OF Triggers on Views

e.g. SQLite, Postgres, Oracle

2. Automatically done by DBMS

e.g. MySQL

INSTEAD OF Triggers on Views

- INSTEAD OF in place of BEFORE or AFTER.
- When an event occurs, the action of the trigger is done instead of the event.
- Translation is put in the action part of the trigger.
- All modifications can be handled.
- No guarantee of correctness

Automatic translation done by DBMS

- Translation to base table is automated by restricting views and modifications.
- No translation needs to be done by SQL programmers.
- Restrictions are significant.
- Adopted by the SQL standard

Example: Modifying a View

```
DROP VIEW if exists LateFeeUsers;  
CREATE VIEW LateFeeUsers AS  
SELECT uID, title, loanDate, overdue  
FROM LOAN  
WHERE overdue =1;
```

```
DELETE FROM LateFeeUsers WHERE uID = 135;
```

SQLite complains that “LateFeeUsers a view and a view can’t be modified !”

Instead of Triggers: Delete

```
Drop trigger LateFeeUsersDelete;  
Create Trigger LateFeeUsersDelete  
instead of delete on LateFeeUsers  
for each row  
begin  
    delete from LOAN  
    where uID = OLD.uID and overdue = 1;  
end;  
  
delete from LateFeeUsers  
where uID = 135;
```

- For each deleted row, the action is taken.
- The deleted row means a tuple that will be logically deleted from the view.
- The OLD variable is the tuple that is asked to be deleted from the view.
- Result: The user 135 is deleted from both LOAN and the view.

Instead of Triggers: Update

```
DROP Trigger LateFeeUsersUpdate;  
CREATE Trigger LateFeeUsersUpdate  
instead of update of title  
ON LateFeeUsers  
for each row  
begin  
    update LOAN  
    set title=New.title  
    where uID = Old.uID and title =  
    Old.title AND overdue = 1;  
end;
```

```
update LateFeeUsers  
set title = 'Bambi II'  
WHERE uID = 24 AND  
title = 'Bambi';
```

Result: The title of uID
24 with 'Bambi II' in
both Loan and the view.

Instead of Trigger: Insert

```
DROP trigger LateFeeUsersInsert;  
CREATE trigger LateFeeUsersInsert  
instead of INSERT ON LateFeeUsers  
for each row  
begin  
    INSERT INTO LOAN VALUES  
    (New.uID, New.title,  
    New.loanDate, New.overdue);  
END;
```

```
INSERT INTO  
LateFeeUsers  
VALUES (888,  
'Gone With the  
Wind', '2000-12-  
25',1) ;
```

Result: The new tuple is inserted in both Loan and the view.

Wrong Translations

- DBMS can't prevent an incorrect trigger (written with logical errors) from being triggered !

Example: Wrong Translations-delete

Drop trigger LateFeeUsersDelete;

Drop trigger WrongTrigger;

Create Trigger WrongTrigger
instead of delete on LateFeeUsers
for each row

begin

delete from LOAN

where uID = OLD.uID and overdue = 1

AND title = 'somebook' ;

end;

delete from LateFeeUsers

where uID =456;

Result: There is no
change in the Loan
Table and the View.

Example: Wrong Translations-update

```
DROP Trigger LateFeeUsersUpdate;  
Drop Trigger WrongTrigger;
```

```
CREATE Trigger WrongTrigger  
instead of update of title ON LateFeeUsers  
for each row  
begin  
  update LOAN  
    set title=New.title  
    where uID = Old.uID and title = Old.title AND  
    loanDate >= datetime('2013-01-01 00:00:00');  
end;
```

```
update LateFeeUsers  
set title = 'Bambi II'  
WHERE uID = 234 AND title = 'Bambi';
```

- Result: There is no change in the view, and a wrong tuple is updated in the Loan table.
- Reason: overdue = 1 is missing and a wrong condition is added in the where clause.

Example: Wrong Translation - insert

- With the LateFeeUsersInsert trigger, do
INSERT INTO LateFeeUsers VALUES
(**111**, 'Modern Database Systems', date('now'), **0**);
- Result: The tuple is inserted into the Loan table but doesn't show up in the view because it doesn't satisfy the condition of the view.
- User should not write an insert to a **view** that will change the base table but doesn't get reflected on the view.

Better LateFeeUsersInsert Trigger

```
DROP trigger LateFeeUsersInsert;  
CREATE trigger LateFeeUsersInsert  
instead of INSERT ON LateFeeUsers  
for each row  
WHEN New.overdue = 1  
begin  
  INSERT INTO LOAN VALUES  
(New.uID, New.title, New loanDate, New.overdue);  
end;
```

```
INSERT INTO LateFeeUsers VALUES(111, 'Streaming Media',  
date('now'), 0);
```

Result: No change in both Loan and the View !

View involving Joins

```
DROP VIEW UsersBeingOverdue;  
CREATE View UsersBeingOverdue as  
  SELECT USER.uID, USER.age, title, overdue  
  FROM USER,LOAN  
  WHERE USER.uID = LOAN.uID AND overdue = 1;
```

View involving Joins: Insert

```
DROP Trigger UsersBeingOverdueInsert;  
CREATE Trigger UsersBeingOverdueInsert  
instead of INSERT ON UsersBeingOverdue  
for each row  
WHEN (New.uID IN (SELECT uID FROM USER)) AND New.overdue = 1  
BEGIN  
    INSERT INTO LOAN VALUES (NEW.uID, New.title, NULL, New.overdue);  
END;  
  
INSERT INTO UsersBeingOverdue VALUES(777, 30,'Lion King', 1); // will not add if  
777 is not in the user relation.  
  
INSERT INTO UsersBeingOverdue  
    SELECT USER.uID, USER.age, 'Java Concepts', 1  
    FROM USER  
    WHERE USER.uID NOT IN (SELECT uID FROM LOAN);
```

View involving Joins: delete

```
DROP Trigger UsersBeingOverdueDelete;  
CREATE Trigger UsersBeingOverdueDelete  
  instead of delete ON UsersBeingOverdue  
  for each row  
  begin  
    delete FROM LOAN  
    WHERE uID = Old.uID AND overdue = 1;  
  end;  
  
DELETE FROM UsersBeingOverdue WHERE uID = 135;
```

View involving Joins: update

```
DROP Trigger UsersBeingOverdueUpdate;  
CREATE Trigger UsersBeingOverdueUpdate  
  instead of update ON UsersBeingOverdue  
  for each row  
  begin  
    update LOAN  
    set overdue = New.overdue  
    WHERE title = NEW.title AND overdue = 1;  
  END;
```

```
UPDATE UsersBeingOverdue  
SET overdue = 0 WHERE title = 'Bambi';
```


View and Constraints

```
CREATE VIEW SeniorUsers AS  
  SELECT uID, uName, age FROM User  
  WHERE age >=60;
```

```
DROP TRIGGER IF EXISTS SeniorUsersInsert;
```

```
CREATE Trigger SeniorUsersInsert  
  instead of insert ON SeniorUsers for each row  
  when New.age >=60  
  begin  
INSERT INTO User VALUES(New.uID, New.uName, New.age, NULL);  
  end;
```

```
INSERT INTO SeniorUsers VALUES(135, 'Kim', 70); //error if 135  
already exists in the User.
```

Ambiguous Translations

- Fundamentally, we could write a trigger for such a modification, but it does not make much sense to write a translation for this type of modification on the view.

Ambiguous Translations

- View involving aggregation

```
DROP VIEW IF EXISTS AvgAge;  
CREATE VIEW AvgAge AS  
SELECT uName, AVG(age) as avg  
FROM User  
GROUP BY uName;
```

- Consider the following query:
update AvgAge set avg = **50** where uName = 'Kim';

Ambiguous Translations

- View using distinct

```
DROP VIEW IF EXISTS UserNames;
```

```
CREATE VIEW UserNames AS
```

```
    SELECT DISTINCT uName FROM User;
```

- Consider the following query:

```
Insert into UserNames values ('Kim');
```

Ambiguous Translations

- View where the sub query references the same table of outer query

```
CREATE VIEW UserWithSameAge AS
```

```
  SELECT * FROM USER U1
```

```
  WHERE EXISTS (SELECT * FROM USER U2
```

```
    WHERE U1.uID <> U2.uID AND U1.age =  
    U2.age);
```

- Consider the following query:
delete from UserWithSameAge where uID =123;

Automatic View Modifications - MySQL

- Restrictions in SQL Standard for “updatable views”
 - SELECT (no DISTINCT) from a single table R (can’t be a join view)
 - Attributes not in SELECT can be NULL or can have a default value.
 - Where clause of a sub query must not involve R in the from clause of outer query (but allow to refer to other tables)
 - No group by or aggregation

Automatic View Modifications

```
DROP VIEW LateFeeUsers;  
CREATE VIEW LateFeeUsers AS  
SELECT uID, title, loanDate, overdue  
FROM LOAN WHERE overdue =1;
```

```
INSERT INTO LateFeeUsers VALUES (1019,  
'Bambi','2000-12-25' , 1); // inserted in both  
Loan and View without any trigger.
```

Example: Wrong Insert

```
INSERT INTO LateFeeUsers VALUES (1018, 'Bambi','2013-12-25' , 0);
```

Result: The tuple is inserted to Loan but the view doesn't reflect the change because overdue = 0.

Solution:

```
DROP VIEW LateFeeUsers;
```

```
CREATE VIEW LateFeeUsers AS
```

```
SELECT uID, title, loanDate, overdue FROM LOAN
```

```
WHERE overdue =1 with check option; // the insert will fail.
```


Example: Not-updatable Views

View with Aggregation

```
CREATE VIEW AvgAge AS  
SELECT uName, AVG(age) as avg  
FROM User  
GROUP BY uName;
```

```
update AvgAge set avg = 50 where uName =  
'Kim'; //error
```

Example: Not-updatable Views

- View with SELECT (distinct)

```
CREATE VIEW UserNames AS
```

```
    SELECT DISTINCT uName FROM User;
```

- Consider the following query:

```
Insert into UserNames values ('Kim'); // error
```

Example: Not-updatable Views

- View with a sub query that refers the same table from outer query.

```
DROP VIEW UserWithSameAge;  
CREATE VIEW UserWithSameAge AS  
  SELECT * FROM USER U1  
  WHERE EXISTS (SELECT * FROM USER U2 WHERE  
U1.uID <> U2.uID AND U1.age = U2.age);
```

Delete from UserWithSameAge where uID =**123**; // error

View involving a subquery that refers to another table

```
DROP VIEW IF EXISTS YoungLoaners;
```

```
Create VIEW YoungLoaners AS
```

```
  SELECT uID, uName, age from User
```

```
  WHERE age < 18 AND User.uID in (select uID from Loan) ;
```

```
delete FROM YoungLoaners WHERE uID = 24;
```

- If the outer query refers a single table, which is required by standard, deletion will be done in the table of the outer query. (Thus, 24 will be removed from User not from Loan)
- If this deletion from the user violates a foreign key constraint, it will be rejected or handled according to a given policy.

View involving a subquery that refers to another table

- Consider

```
INSERT INTO YoungLoaners VALUES (187, 'Wu', 60);
```

- Result: Wu will be inserted to User in the outer query but will not be shown in the view due to his age. This insertion to the user through the view YoungLoaners is not desirable.
- Solution: check with option. This insertion will be failed by the check option.

Join Views

- SQL Standard doesn't allow to update Join Views.
- MySQL allows insert and update operations on join views, but not delete operations.

Why Join Views are not updatable by standard ?

```
CREATE VIEW MovieProd AS
```

```
  SELECT title, year, name
```

```
  FROM Movies, MovieExec
```

```
  WHERE producerC# = cert#;
```

```
Insert into MovieProd Values ('Bambi', 1969,  
'Max',);
```

Translation:

```
Insert into Movies Values ('Bambi', 1969, null);
```

```
Insert into MovieExec ('Max', null);
```

MySQL: Join Views

Drop View if EXISTS BambiUsers;

Create view BambiUsers(uID, loanID, uName, overdue) AS

Select User.uID, Loan.uID, User.uName, overdue

From **User, Loan**

Where User.uID = Loan.uID and title = 'Bambi' ;

- update BambiUsers set overdue = **0** WHERE overdue = **1**;

Result: The system will modify the base table containing the attribute to be updated. (Loan in this case.)

- update BambiUsers set loanID = **111** WHERE loanID = **1011**;

Result: Without check option and without foreign key constraint, the uID in Loan is updated, but the corresponding uID in User is not changed.

With check option, it will CHECK OPTION fail if 111 doesn't exist in User.

Without check option and with foreign key constraint, it will fail due to the foreign key violation if 111 doesn't exist in User.

Indexes

- Index on an attribute A: A data structure that makes it faster to find those tuples that have specific column values
- A data structure of (key, value)
 - Key x: A value of attribute A \rightarrow index key
 - Value: set of locations of the tuples that have x for the value of A.
- Index key can be any attribute or set of attributes
- Stored in database
- Underlying data structures
 - B tree/B+ tree
 - Hash table

Motivation

```
SELECT uID  
FROM User, Loan  
WHERE User.uID = Loan.uID and title = 'Bambi' ;
```

If we have an index on title of Loan, first get the tuples containing Bambi, and test uIDs of the tuples for equality.

Defining Indexes

```
CREATE INDEX TitleIndex ON Loan (title);
```

To serve the query in the previous slide, the DBMS can quickly determine the position of tuples that contain 'Babmi' without having to look at all the data.

A single index on multiple attributes

- Suppose title and year form a primary key for Movie. Then consider

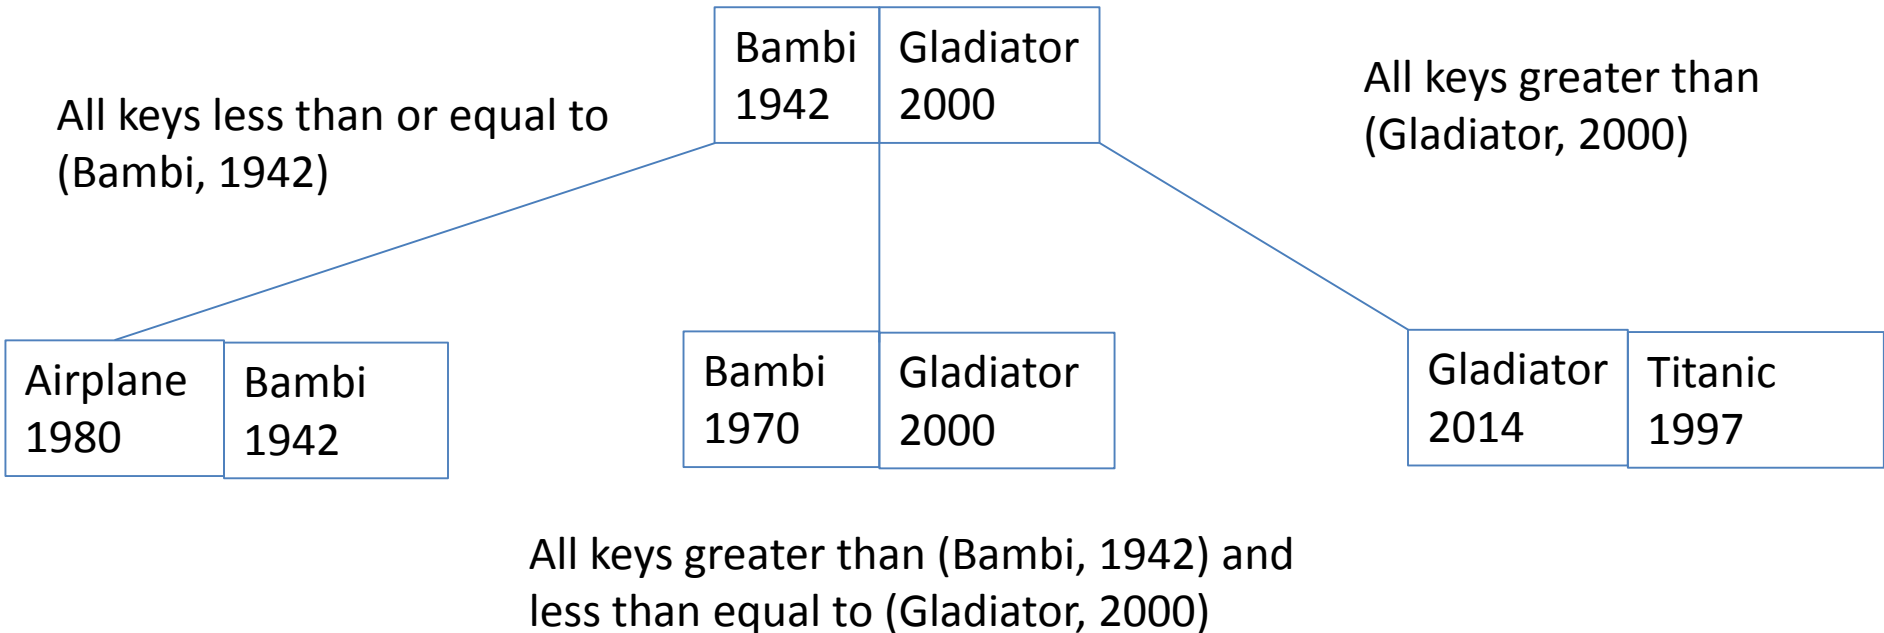
```
SELECT *  
FROM Movie  
Where title = 'Star Wars' and year = 1990;
```

- **With** `CREATE INDEX KeyIndex Movie(title, year)`
The index will find only one tuple.
- **With** `CREATE INDEX YearIndex Movie(year)`
First find all movies made on 1990 using the index and check through them for the given title.
- `(title, year)` vs. `(year, title)`

Composite Index Key

- The ordering of the composite index key values is lexicographic ordering. That is, $(a_1, a_2) < (b_1, b_2)$ if either $a_1 < b_1$ or $(a_1 = b_1 \text{ and } a_2 < b_2)$
- The order of the individual attributes in the composite key is important.

Composite Index Key



The above index may be used to search for (title, year) or (title), but not for (year).

Selection of Indexes

- Benefits: faster query execution
- Costs: space, index creation, maintenance
 - Maintenance is the most expensive overhead: Insertions, deletions and updates to that relation become more complex and time-consuming.

Selection of Indexes

- The most useful index on a relation is an index on its key (key index)
 - Queries involving the key are very common.
 - At most one disk page needs to be retrieved because a key is associated with at most one tuple.
- If the index key is not the key of a relation
 - The index key needs to be almost a key
(e.g. title instead of the key (title, year) for Movies)
 - or
 - The tuples needs clustered on the index key
in order to minimize the time to retrieve tuples from the disk.

Example: Selection of Index

- **With** `CREATE INDEX YearIndex Movie(year)`

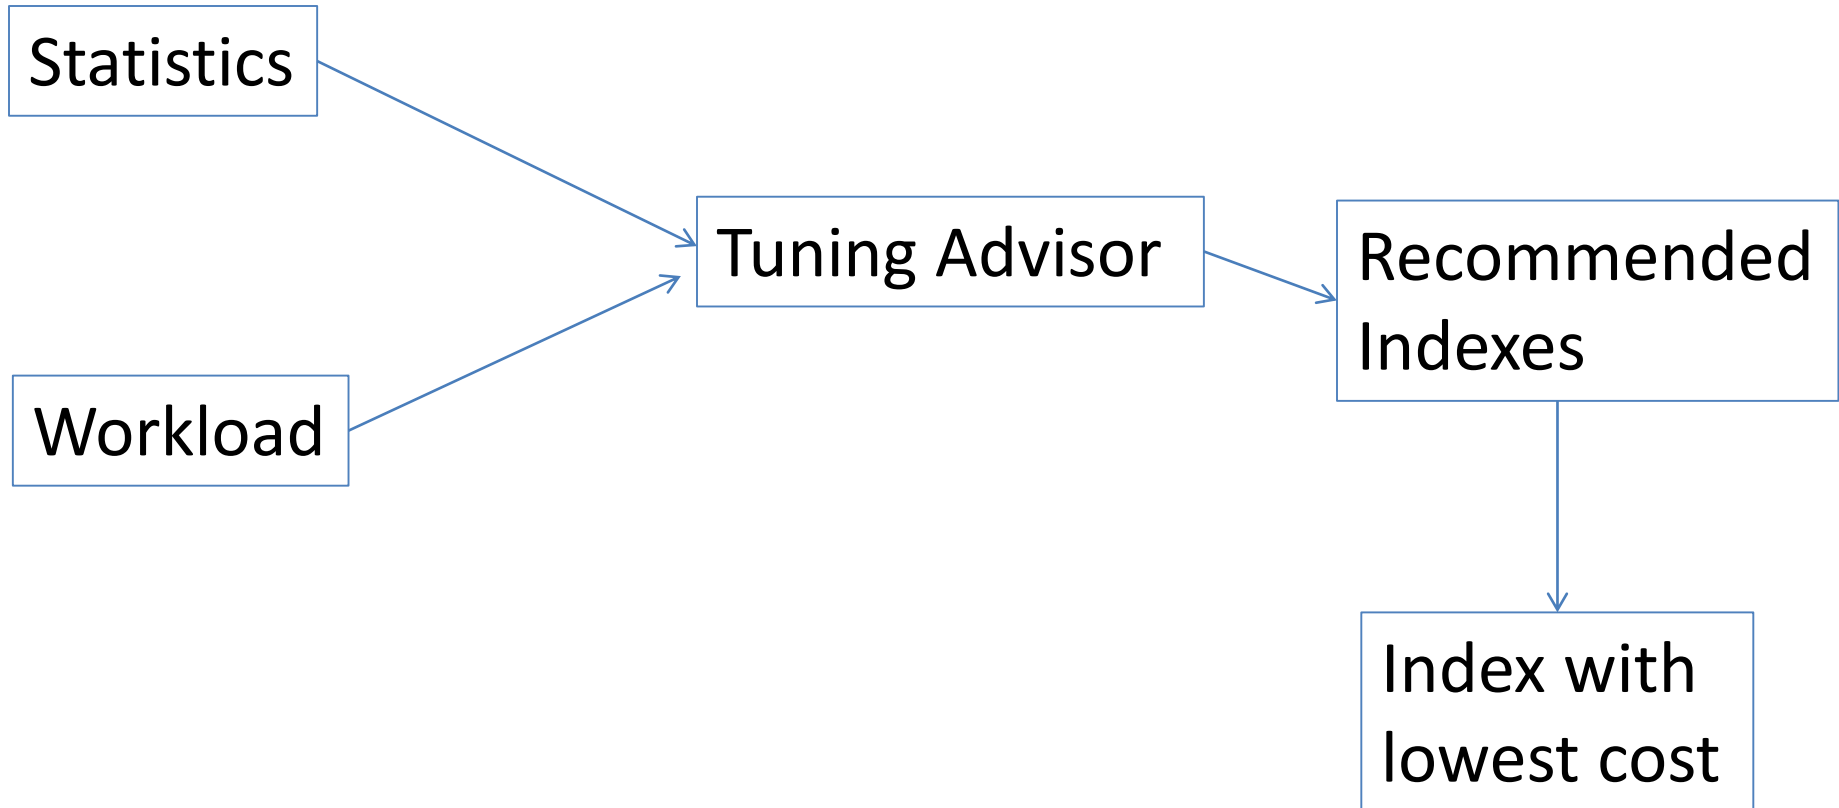
- **Consider**

```
select * from Movies where year= 1990;
```

- **What if Movies are not clustered by year ?**

There will be little gain from the index on year. Imagine a scenario where movie tuples made on 1990 are spread all over the places on the disk.

Index Creation



Example: Index Creation

`StarIn(movieTitle, movieYear, starName)`

Assumptions:

- `StarIn` occupies 10 disk pages
- A star has appeared in 3 movies
- A movie has 3 stars
- One disk access is assumed to read a page of the index every time we use that index.
- Modification on the index needs two disk accesses (one to read the page and another to write it back to disk)

Example: Index Creation

Q1:

```
SELECT movieTitle, movieYear  
FROM StarsIn  
WHERE starName = s;
```

No Index	Star Index	Movie Index	Both Index
10	4 (1 index access + 3 tuple accesses)	10	4 (movie index doesn't help)

Example: Index Creation

Q2 :

```
SELECT starName
```

```
FROM StarsIn
```

```
WHERE movieTitle=t AND movieYear = y;
```

No Index	Star Index	Movie Index	Both Index
10	10	4	4 (star index doesn't help)

Example: Index Creation

I :

```
INSERT INTO StarIn VALUES (t, y, s) ;
```

No Index	Star Index	Movie Index	Both Index
2	4 (2 for modifying data and 2 for modifying index)	4	6

Example: Index Creation

	No Index	Star Index	Movie Index	Both Index
Average	$8P_1+8P_2+2$	$6P_2+4$	$6P_1+4$	$6-2P_1-2P_2$

P_1 : Fraction of time doing Q1

P_2 : Fraction of time doing Q2

$I: 1 - P_1 - P_2$

If insertions are dominant: No Index

If queries are dominant: Both Index

Materialized Views

- Stored in database
- Benefits of using Virtual Views + faster query performance.
- In principle, we need to re-compute a materialized view every time one of its base tables changes
- Make changes to the materialized view incremental.

Incremental Changes to Materialized Views

```
CREATE MATERIALIZED VIEW MovieProd AS  
  SELECT title, year, name  
  FROM Movies, MovieExec  
  WHERE producerC# = cert#;
```

```
Movies(title, year, producerC#)  
MovieExec(name, cert#)
```

Example: Incremental Changes to Materialized Views

```
INSERT INTO Movies VALUES ('Kill Bill', 2003, 23456);
```



```
SELECT name
```

```
FROM MovieExec
```

where cert# = 23456; // only one name, say Smith,
returns because cert# is the key

```
INSERT INTO MovieProd VALUES ('Kill Bill', 2003,  
'Smith');
```

Example: Incremental Changes to Materialized Views

```
DELETE FROM Movies WHERE title = 'Bambi'  
and year = 1994;
```



```
DELETE FROM MovieProd  
WHERE title = 'Bambi' and year = 1994;
```

Example: Incremental Changes to Materialized Views

```
INSERT INTO MovieExec VALUES ('Max', 34567);
```



```
INSERT INTO MovieProd  
  SELECT title, year, 'Max'  
  FROM Movies  
  WHERE prodecerc# = 34567;
```

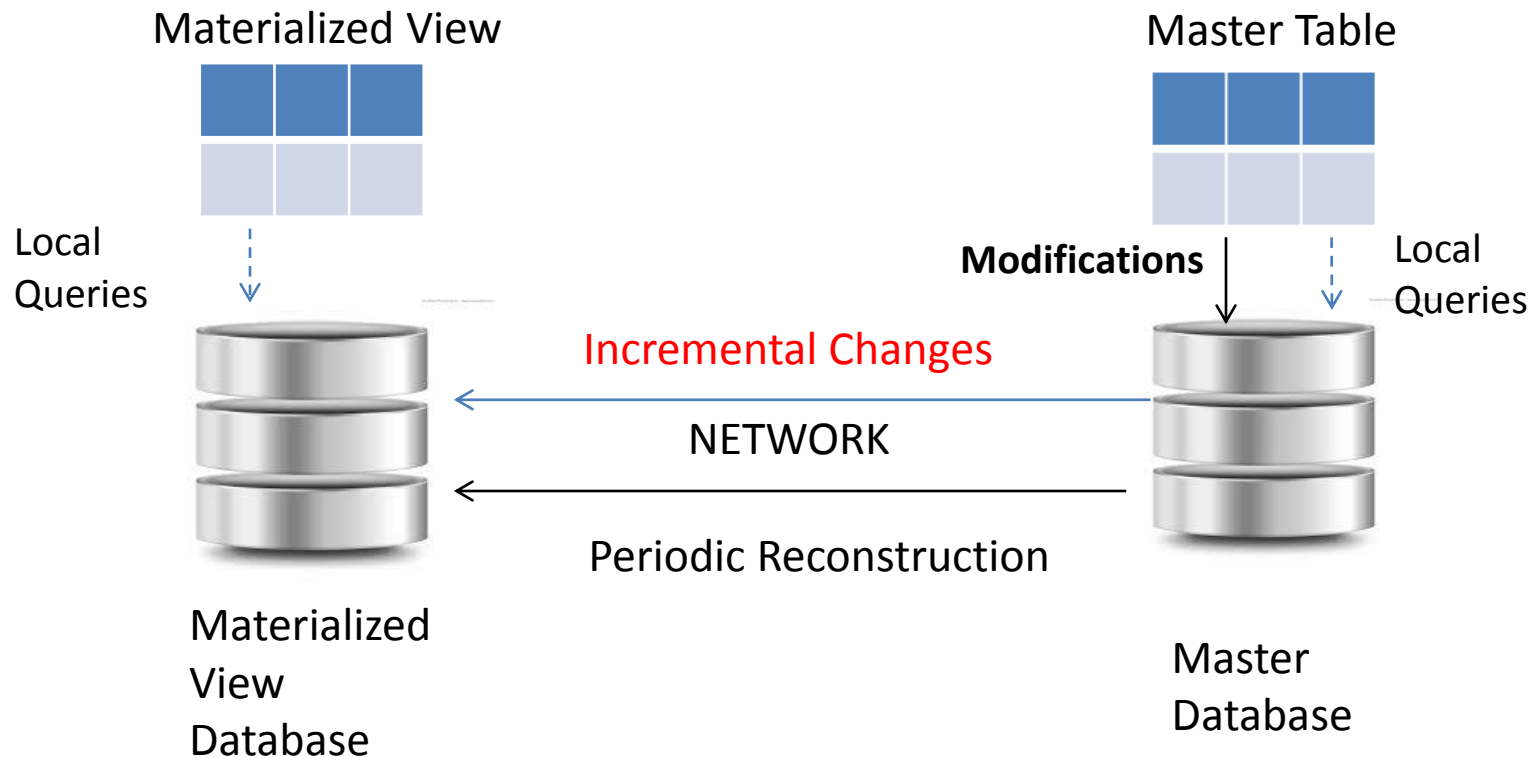
Example: Incremental Changes to Materialized Views

```
DELETE FROM MovieExec WHERE cert#=45678;
```



```
DELETE FROM MovieProd  
WHERE (title, year) IN  
      (SELECT title, year FROM Movies  
       WHERE producerC# = 45678);
```

Periodic Maintenance of Materialized Views



Rewriting Queries to Use Materialized Views

Materialized View

Query to be rewritten

SELECT Lv (3) Lq that come from Rv are on Lv
FROM Rv (1) All Rv appear in Rq
WHERE Cv (2) Cq = Cv AND C
If C refers attributes of Rv,
they should be
on Lv

SELECT Lq
FROM Rq
WHERE Cq

SELECT title, year, name
FROM Movies, MovieExec
WHERE producer#=cert#;

SELECT starName
FROM StarIn, Movies, MovieExec
WHERE movieTitle = title AND movieYear=year
AND producer#=cert# AND name = 'Max';

Movies(title,year,length,genre,studioName,producerC#)
StarIn(movieTitle, movieYear, starName)
MovieExec(name, address, cert#, netWorth)

Rewriting Queries to Use Materialized Views

Materialized View

```
SELECT title, year, name  
FROM Movies, MovieExec  
WHERE producer#=cert#;
```

Query to be rewritten

```
SELECT starName  
FROM StarIn, Movies, MovieExec  
WHERE movieTitle = title AND movieYear=year  
AND producer#=cert# AND name = 'Max';
```

Rewritten Query

```
SELECT starName  
FROM StarIn, MovieProd  
WHERE movieTitle=title AND movieYear = year AND name = 'Max';
```

The same result but faster!